

template

March 21, 2024

1 Recipe Data Analytics

Name(s): Daniel Warren

Website Link: <https://danielpwarren.github.io/Recipe-Data-Analytics>

```
[ ]: import pandas as pd
import numpy as np
from pathlib import Path
import ast # To convert strings that resemble lists into actual lists
from itertools import chain
from sklearn.feature_extraction.text import TfidfVectorizer
import math

import plotly.express as px

pd.options.plotting.backend = "plotly"
```

1.1 Step 1: Introduction

1.1.1 Brainstormed Questions

- How do the ingredients and nutrition contribute to the rating of a recipe?
- How do the number of steps and expected time contribute to the rating?
- Are certain food categories (such as desert or dinner) more likely to receive higher ratings?
- What terms in the description or tags correlate to higher ratings?
- What terms in the reviews correlate to higher ratings?

1.1.2 Selected Question

How do the number of steps and expected time contribute to the rating?

1.1.3 Loading the data:

```
[ ]: recipes_raw = pd.read_csv(Path("data") / Path("RAW_recipes.csv"))
interactions_raw = pd.read_csv(Path("data") / Path("RAW_interactions.csv"))
```

```
[ ]: interactions_raw.columns
```

```
[ ]: Index(['user_id', 'recipe_id', 'date', 'rating', 'review'], dtype='object')
```

1.2 Step 2: Data Cleaning and Exploratory Data Analysis

1.2.1 Data Cleaning

In this step I will:

- combine the recipes and interactions dataframes into one dataframe.
- separate `tags`, `steps`, and `ingredients` into arrays
- replace 0 star ratings with NaN and calculate an average rating for each recipe and store it in `avg_rating`
- change the datatypes of rows where applicable

```
[ ]: interactions_raw["rating"] = interactions_raw["rating"].replace(0, np.nan)

recipes_raw["tags"] = recipes_raw["tags"].apply(ast.literal_eval)
recipes_raw["steps"] = recipes_raw["steps"].apply(ast.literal_eval)
recipes_raw["ingredients"] = recipes_raw["ingredients"].apply(ast.literal_eval)
recipes_raw["nutrition"] = recipes_raw["nutrition"].apply(ast.literal_eval)

merged = recipes_raw.merge(
    interactions_raw, left_on="id", right_on="recipe_id", how="left"
)
merged.head()

avg_rating = merged.groupby("id")["rating"].mean()
avg_rating
merged = merged.set_index("id")
merged["avg_rating"] = avg_rating
recipes = merged.reset_index()

recipes["date"] = pd.to_datetime(recipes["date"])
recipes["submitted"] = pd.to_datetime(recipes["submitted"])
recipes.drop(columns=["recipe_id"], inplace=True)

nutrient_names = [
    "calories",
    "total fat (PDV)",
    "sugar (PDV)",
    "sodium (PDV)",
    "protein (PDV)",
    "saturated fat (PDV)",
    "carbohydrates (PDV)",
]
for index, nutrient in enumerate(nutrient_names):
    recipes[nutrient] = recipes["nutrition"].apply(lambda x: x[index])
    recipes[nutrient] = pd.to_numeric(recipes[nutrient])
```

```
print(recipes.dtypes)
```

```
id                int64
name              object
minutes           int64
contributor_id    int64
submitted         datetime64[ns]
tags              object
nutrition         object
n_steps           int64
steps             object
description       object
ingredients       object
n_ingredients     int64
user_id           float64
date              datetime64[ns]
rating            float64
review           object
avg_rating        float64
calories          float64
total fat (PDV)   float64
sugar (PDV)       float64
sodium (PDV)      float64
protein (PDV)     float64
saturated fat (PDV) float64
carbohydrates (PDV) float64
dtype: object
```

Note: user_id, rating, and avg_rating here all need to be floats as the columns contain NaN values that can't be represented as integers.

```
[ ]: for col in recipes.columns:
      print(f"{col}: {type(recipes[col][0])}")
```

```
id: <class 'numpy.int64'>
name: <class 'str'>
minutes: <class 'numpy.int64'>
contributor_id: <class 'numpy.int64'>
submitted: <class 'pandas._libs.tslibs.timestamps.Timestamp'>
tags: <class 'list'>
nutrition: <class 'list'>
n_steps: <class 'numpy.int64'>
steps: <class 'list'>
description: <class 'str'>
ingredients: <class 'list'>
n_ingredients: <class 'numpy.int64'>
user_id: <class 'numpy.float64'>
date: <class 'pandas._libs.tslibs.timestamps.Timestamp'>
```

```

rating: <class 'numpy.float64'>
review: <class 'str'>
avg_rating: <class 'numpy.float64'>
calories: <class 'numpy.float64'>
total fat (PDV): <class 'numpy.float64'>
sugar (PDV): <class 'numpy.float64'>
sodium (PDV): <class 'numpy.float64'>
protein (PDV): <class 'numpy.float64'>
saturated fat (PDV): <class 'numpy.float64'>
carbohydrates (PDV): <class 'numpy.float64'>

```

1.2.2 Univariate Analysis

In this step I will observe the distributions of relevant columns in the dataframe

```

[ ]: fig1 = px.bar(recipes["rating"].value_counts(), y="rating")
fig1.update_layout(title="Distribution of ratings")
fig1.update_xaxes(title="rating")
fig1.update_yaxes(title="count")
fig1.show()

[ ]: fig2 = px.bar(recipes[recipes["n_steps"] <= 40]["n_steps"].value_counts(),
    ↪y="n_steps")
fig2.update_layout(title="Distribution of number of steps")
fig2.update_xaxes(title="steps")
fig2.update_yaxes(title="count")
fig2.show()

[ ]: bins = np.arange(0, 12 * 60 + 15, 15)
bin_indices = np.digitize(recipes[recipes["minutes"] <= 12 * 60]["minutes"],
    ↪bins) - 1
result = np.bincount(bin_indices)
df = pd.DataFrame({"minute_bin": bins, "count": result})
fig3 = px.bar(
    df, x="minute_bin", y="count", title="Distribution of Minutes in 15-Minute_
    ↪Bins"
)
fig3.update_xaxes(title="minutes")
fig3.update_yaxes(title="count")
fig3.show()

[ ]: all_tags = list(chain.from_iterable(recipes["tags"]))
tag_counts = pd.Series(all_tags).value_counts()

percentile_threshold = tag_counts.quantile(0.90)
truncated_counts = tag_counts[tag_counts >= percentile_threshold]
truncated_counts

```

```
fig4 = px.bar(truncated_counts)
fig4.update_layout(title="Distribution of most common tags (top 10%)")
fig4.update_xaxes(title="tag")
fig4.update_yaxes(title="count")
fig4.show()
```

1.2.3 Bivariate Analysis

```
[ ]: fig4 = px.scatter(recipes, x="avg_rating", y="n_steps")
fig4.show()
```

```
[ ]: recipes["truncated_minutes"] = recipes[recipes["minutes"] <= 12 * 60]["minutes"]
fig5 = px.scatter(recipes, x="avg_rating", y="truncated_minutes")
fig5.show()
```

1.2.4 Interesting Aggregates

```
[ ]: recipes_df = recipes.copy()

recipes_df = recipes_df[recipes_df["minutes"] <= 200] # get rid of outliers
pivot_table = recipes_df.pivot_table(
    values="n_steps", index="n_ingredients", aggfunc=["mean", "median", "min", "max"]
)

pivot_table
```

```
[ ]:
```

	mean	median	min	max
	n_steps	n_steps	n_steps	n_steps
n_ingredients				
1	6.750000	7	2	20
2	5.969932	5	1	55
3	5.539765	5	1	69
4	6.493830	5	1	55
5	7.329812	6	1	80
6	7.740900	7	1	86
7	8.297483	8	1	52
8	9.052886	8	1	67
9	9.759984	9	1	87
10	10.698326	10	1	57
11	11.187419	10	1	65
12	11.782152	11	1	57
13	12.549089	11	1	68
14	13.540655	12	1	88
15	14.356427	13	2	62
16	15.142857	14	2	77
17	15.656934	15	1	55

18	16.410724	15	3	65
19	16.675875	16	3	59
20	17.226623	17	2	76
21	18.733615	17	3	48
22	21.721402	27	4	59
23	17.753676	17	5	59
24	17.673203	17	4	44
25	19.140000	18	6	45
26	18.948454	21	5	49
27	18.088235	19	7	34
28	23.942857	23	5	76
29	18.551724	17	15	33
30	27.906250	23	6	62
31	18.916667	20	10	29
32	37.000000	40	28	40
33	6.000000	6	6	6

1.3 Step 3: Assessment of Missingness

1.3.1 NMAR Analysis

We will consider the `review` column. This column is the most likely to have NMAR missing data, since recipes with medium reviews could be correlated with missing review text.

1.3.2 Missingness Dependency

We will further analyze the `rating` column, against the `n_ingredients` and `minutes` columns

```
[ ]: recipes["rating"].isna().mean()
```

```
[ ]: 0.06413882241531552
```

```
[ ]: def dependency_test(df, col):
    print(f"Testing dependency of rating on {col}")
    observed_diff = abs(
        df[col][df["rating"].isna()].mean() - df[col][~df["rating"].isna()].
        ↪mean()
    )
    sim_diff = []
    shuffled = df.copy()
    for _ in range(1000):
        shuffled[col] = np.random.permutation(shuffled[col])
        sim_diff.append(
            abs(
                shuffled[col][shuffled["rating"].isna()].mean()
                - shuffled[col][~shuffled["rating"].isna()].mean()
            )
        )
```

```

print(observed_diff)
print(f"{sim_diff[:5]}...{sim_diff[-5:]}")
print((sim_diff >= observed_diff).mean())

```

```

dependency_test(recipes, "n_steps")
dependency_test(recipes, "n_ingredients")
dependency_test(recipes, "minutes")

```

Testing dependency of rating on n_steps

1.3386412335909217

[0.003422808128130228, 0.029716887844099205, 0.026874284631562162,
0.03580511620487847, 0.0012197906384159296]...[0.03516553048205928,
0.028414347852283584, 0.0817131580873518, 0.09159120425091771,
0.04407203406741189]

0.0

Testing dependency of rating on n_ingredients

0.1607379066254797

[0.022823195824098974, 0.03227485150578424, 0.00022450028443010694,
0.0763125912131315, 0.03253650174005962]...[0.008610179761413761,
0.059825492580415585, 0.07062738478805564, 0.004963250578450484,
0.05494461212576773]

0.0

Testing dependency of rating on minutes

51.45237039852127

[3.538001943248858, 10.787350786021463, 5.28527907291506, 29.27208882634673,
3.7211366552165543]...[21.879845332454295, 20.743497472852155,
16.779131967567665, 62.03786943660768, 9.079514775929198]

0.114

1.4 Step 4: Hypothesis Testing

The hypotheses I plan to test are:

H_0 : The number of steps and the expected time to complete a task have no effect on the task's rating. Any observed differences in ratings due to changes in the number of steps and expected time are purely due to chance.

H_1 : The number of steps and the expected time to complete a task both negatively affect the task's rating. Tasks that have a higher number of steps and longer expected times to complete will have lower ratings, not due to chance.

```

[ ]: def permutation_test_correlation(df, col, rating_col="avg_rating"):
    clean_df = df.dropna(subset=[col, rating_col])

    print(f"Testing dependency of {rating_col} on {col}")
    observed_corr = clean_df[[col, rating_col]].corr().iloc[0, 1]

    sim_corr = []

```

```

for _ in range(1000):

    shuffled_rating = np.random.permutation(clean_df[rating_col].values)

    sim_corr_value = np.corrcoef(clean_df[col].values, shuffled_rating)[0,1]
    ↪1]

    if not np.isnan(sim_corr_value):
        sim_corr.append(sim_corr_value)

    p_value = np.mean(np.abs(sim_corr) >= np.abs(observed_corr))
    print(f"Observed correlation: {observed_corr}")
    print(f"P-value: {p_value}")

permutation_test_correlation(recipes, "n_steps")
permutation_test_correlation(recipes, "minutes")

```

Testing dependency of avg_rating on n_steps
 Observed correlation: -0.0018214513463115766
 P-value: 0.387
 Testing dependency of avg_rating on minutes
 Observed correlation: 0.00196393037826293
 P-value: 0.259

1.5 Step 5: Framing a Prediction Problem

I plan to predict the calories of each recipe.

The features I'll use for this are as follows - total fat (PDV) - protein (PDV)

1.6 Step 6: Baseline Model

We will begin with a very rudimentary baseline model. We'll use a simple LinearRegression model and train it to predict calories based only on the recipes total fat and sugar.

```

[ ]: from sklearn.model_selection import train_test_split
    from sklearn.pipeline import Pipeline
    from sklearn.linear_model import LinearRegression
    from sklearn.metrics import mean_squared_error, r2_score

    data = recipes[["total fat (PDV)", "protein (PDV)", "calories"]].
    ↪drop_duplicates()

    X = data.drop(["calories"], axis=1)
    y = data["calories"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
    )

```



```

baseline_pipeline = Pipeline(
    [
        ("regressor", LinearRegression()),
    ]
)

baseline_pipeline.fit(X_train, y_train)
y_pred = baseline_pipeline.predict(X_test)

mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

(mse, r2)

```

```
[ ]: (156606.05998095873, 0.6582566526599032)
```

1.7 Step 7: Final Model

Here I create new features by dividing calories against the gram measurements of protein, carbs, and fat.

```

[ ]: from sklearn.model_selection import GridSearchCV
from sklearn.pipeline import Pipeline
from sklearn.compose import ColumnTransformer
from sklearn.linear_model import Lasso
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import Binarizer, PolynomialFeatures

feature_engineering = ColumnTransformer(
    [
        # Fitting a polynomial maps to the total fat should increase the
        # accuracy of the model as it carries more calories
        ("poly_fat", PolynomialFeatures(degree=2), ["total fat (PDV)"]),
        # As well, since the fat seems to be the most important feature, we can
        # binarize it to see if it has a significant impact on the model
        ("high_fat", Binarizer(threshold=21.0), ["total fat (PDV)"]),
    ],
    remainder="passthrough",
)

final_model_pipeline = Pipeline(
    [
        ("feature_engineering", feature_engineering),
        ("regressor", Lasso()),
    ]
)

```

```

)

param_grid = {
    "regressor__alpha": [0.0001, 0.001, 0.01, 0.1],
    "regressor__max_iter": [1000, 10000, 100000],
    "regressor__tol": [0.0001, 0.001, 0.01, 0.1],
    "regressor__fit_intercept": [True, False],
}

grid_search = GridSearchCV(final_model_pipeline, param_grid, cv=5, scoring="r2")
grid_search.fit(X_train, y_train)

best_params = grid_search.best_params_
best_score = grid_search.best_score_

print("Best Parameters:", best_params)
print("Best R-squared:", best_score)

y_pred = grid_search.predict(X_test)
new_mse = mean_squared_error(y_test, y_pred)
new_r2 = r2_score(y_test, y_pred)
print("Test MSE:", mse)
print("Test R-squared:", r2)

```

Best Parameters: {'regressor__alpha': 0.1, 'regressor__fit_intercept': False, 'regressor__max_iter': 1000, 'regressor__tol': 0.01}

Best R-squared: 0.7983434554034269

Test MSE: 156606.05998095873

Test R-squared: 0.6582566526599032

```

[ ]: feature_engineering = ColumnTransformer(
    [
        ("poly_fat", PolynomialFeatures(degree=2), ["total fat (PDV)"]),
        ("high_fat", Binarizer(threshold=33.6), ["total fat (PDV)"]),
    ],
    remainder="passthrough",
)

final_model_pipeline = Pipeline(
    [
        ("feature_engineering", feature_engineering),
        ("regressor", Lasso(alpha=0.1, max_iter=1000, tol=0.01,
fit_intercept=False)),
    ]
)

final_model_pipeline.fit(X_train, y_train)

```

```

y_pred = final_model_pipeline.predict(X_test)
new_mse = mean_squared_error(y_test, y_pred)
new_r2 = r2_score(y_test, y_pred)

print("Test MSE:", new_mse)
print("Test R-squared:", new_r2)
print("MSE Improvement:", mse - new_mse)
print("R-squared Improvement:", new_r2 - r2)

```

```

Test MSE: 155977.12096335526
Test R-squared: 0.659629113758629
MSE Improvement: 628.9390176034649
R-squared Improvement: 0.0013724610987257968

```

1.8 Step 8: Fairness Analysis

The fairness analysis for this model is an interesting one as the model I ended up with is very simple. Still, we will create a group using the mean **protein** (PDV) as this seems to be a less decisive element in our model so it will be interesting to see if the models RMSE for both sides of the value are the same.

H_0 : The model is fair, its mean RMSE for values lower than the mean protein in the dataset is close to the RMSE of values greater than the mean protein

H_1 : The model is not fair and there is a bias towards either lower or higher protein values

```

[ ]: from sklearn.metrics import mean_squared_error
import numpy as np

protein_median = data["protein (PDV)"].mean()

high_protein_test = X_test[
    data.loc[X_test.index, "protein (PDV)"] > protein_median
]
low_protein_test = X_test[
    data.loc[X_test.index, "protein (PDV)"] <= protein_median
]

y_test_high_protein = y_test[high_protein_test.index]
y_test_low_protein = y_test[low_protein_test.index]

y_pred_high_protein = final_model_pipeline.predict(high_protein_test)
y_pred_low_protein = final_model_pipeline.predict(low_protein_test)

rmse_high_calorie = np.sqrt(
    mean_squared_error(y_test_high_protein, y_pred_high_protein)
)

```

```

rmse_low_calorie = np.sqrt(mean_squared_error(y_test_low_protein,
↪y_pred_low_protein))

observed_diff = rmse_high_calorie - rmse_low_calorie

rmse_high_calorie, rmse_low_calorie, observed_diff

```

```
[ ]: (625.5672862708465, 168.57409034869306, 456.9931959221534)
```

```

[ ]: def permutation_test(
    y_true_high, y_pred_high, y_true_low, y_pred_low, n_permutations=1000
):
    y_concat = np.concatenate([y_true_high, y_true_low])

    observed_rmse_diff = np.sqrt(
        mean_squared_error(y_true_high, y_pred_high)
    ) - np.sqrt(mean_squared_error(y_true_low, y_pred_low))

    permuted_diffs = []

    for _ in range(n_permutations):
        np.random.shuffle(y_concat)

        y_perm_high = y_concat[: len(y_true_high)]
        y_perm_low = y_concat[len(y_true_high) :]

        rmse_perm_high = np.sqrt(mean_squared_error(y_perm_high, y_pred_high))
        rmse_perm_low = np.sqrt(mean_squared_error(y_perm_low, y_pred_low))

        permuted_diffs.append(rmse_perm_high - rmse_perm_low)

    p_value = np.mean(np.abs(permuted_diffs) >= np.abs(observed_rmse_diff))

    return p_value

p_value = permutation_test(
    y_test_high_protein,
    y_pred_high_protein,
    y_test_low_protein,
    y_pred_low_protein,
    n_permutations=10000,
)
print(f"P-value from permutation test: {p_value}")

```

P-value from permutation test: 0.0933